

OPAM: a Package Management Systems for OCaml
Version 1.0.0 Roadmap

THIS DOCUMENT IS A DRAFT

Thomas GAZAGNAIRE
thomas.gazagnaire@ocamlpro.com

June 15, 2012

Contents

1	Milestone 1: Foundations	3
1.1	Client state	3
1.1.1	Global state	4
1.1.2	Compiler-specific state	4
1.1.3	Repository-specific state	5
1.2	File syntax	5
1.2.1	List of packages	5
1.2.2	Index of packages	6
1.2.3	General syntax	6
1.2.4	Global configuration file	7
1.2.5	Package specification files: <code>.opam</code>	7
1.2.6	Package configuration files: <code>.config</code>	9
1.2.7	Package installation files: <code>.install</code>	10
1.2.8	Substitution files	10
1.3	Client commands	10
1.3.1	Creating a fresh client state	10
1.3.2	Listing packages	11
1.3.3	Getting package info	11
1.3.4	Installing a package	12
1.3.5	Updating index files	13
1.3.6	Upgrading installed packages	13
1.3.7	Getting package configuration	13
1.3.8	Uploading packages	14
1.3.9	Removing packages	15
1.3.10	Managing OPAM repository	15
1.3.11	Dependency solver	15
1.4	OPAM repository scripts	15
1.4.1	Init script	16
1.4.2	Update script	16
1.4.3	Download script	16
1.4.4	Upload script	16
1.4.5	Error handling	16
1.4.6	RSYNC repository	16
2	Milestone 2: OPAM server	17
2.1	Server state	17
2.2	Binary Protocol	17
2.3	Server authentication	18
2.4	Scripts	19
3	Milestone 3: Multiple compiler versions	19
3.1	Compiler Description Files	19
3.2	Switching compiler version	21
3.3	OPAM repositories	21
4	Milestone 4: OPAM and git	21
4.1	OPAM repository under git	21
4.2	Package repository under git	21

5 Milestone 9: Parallel Build	21
6 Milestone 8: Version Pinning	21

Overview

This document specifies the design of a package management system for OCaml (OPAM). For the first version of OPAM, we have tried to consider the simplest design choices, even if these choices restrict user possibilities (but we hope not too much). Our goal is to propose a system that we can build in a few months. Some of the design choices might evolve to more complex tasks later, if needed.

A package management system has typically two kinds of users: *end-users* who install and use packages for their own projects; and *packagers*, who create and upload packages. End-users want to install on their machine a consistent collection of *packages* – a package being a collection of OCaml libraries and/or programs. Packagers want to take a collection of their own libraries and programs and make them available to other developpers.

This document describes the fonctional requirements for both kinds of users.

Conventions

In this document, `$home`, `$opam`, `$package` and `$path` are assumed to be defined as follows:

- `$home` refers to the end-user home path, typically `/home/thomas/` on linux, `/Users/thomas/` on OSX `C:\Documents and Settings\thomas\` on Windows.
- `$opam` refers to the filesystem subtree containing the client state. Default directory is `$home/.opam`.
- `$package` refers to a path in the packager filesystem, where lives the collection of libraries and programs he wants to package.
- `$path` refers to a list of paths in the packager filesystem, where lives the collection of programs (`ocamlc`, `ocamldep`, `ocamlopt`, `ocamlbuild`, ...).

User variables are written in capital letters, prefixed by `$`. For instance package names will be written `$NAME`, package versions `$VERSION`, and the version of the ocaml compiler currently installed `$OVERSION`.

1 Milestone 1: Foundations

The first milestone of OPAM focuses on providing a limited set of features, dedicated to package management of OCaml packages. OPAM rely on external tools to compile and provide full configuration options to the build tools. The goal for this first milestone is to be as much as possible compatible with any existing build system (including `ocamlfind` and `oasis`) modulo few modifications.

1.1 Client state

The client state is stored on the filesystem, under `$opam`:

1.1.1 Global state

- `$opam/config` is the main configuration file. It defines the OPAM version, the repository addresses and the current compiler version. The file format is described in §1.2.4.
- `$opam/opam/$NAME.$VERSION.opam` is the OPAM specification for the package `$NAME` with version `$VERSION` (which might not be installed). The format of OPAM files is described in §1.2.5.
- `$opam/descr/$NAME.$VERSION` contains the description for the version `$VERSION` of package `$NAME` (which might not be installed). The first line of this file is the package synopsis.
- `$opam/archives/$NAME.$VERSION.tar.gz` contains the source archives for the version `$VERSION` of package `$NAME`.
- `$opam/repo/index` contains the location of packages (ie. in which repositories they can be found and the priority between repositories). The file format is described in 1.2.2.

1.1.2 Compiler-specific state

All the configurations files, libraries and binaries related to the specific `$OVERSION` of the OCaml compiler are stored in `$opam/$OVERSION`.

- `$opam/$OVERSION/installed` is the list of installed packages for the compiler version `$OVERSION`. The file format is described in §1.2.1.
- `$opam/$OVERSION/config/$NAME.config` is a platform-specific configuration file of for the installed package `$NAME` with the compiler version `$OVERSION`. The file format is described in §1.2.6. `$opam/$OVERSION/config/` can be shortened to `$config/` for more readability.
- `$opam/$OVERSION/install/$NAME.install` is a platform-specific package installation file for the installed package `$NAME` with the compiler version `$OVERSION`. The file format is described in §1.2.7. `$opam/$OVERSION/install` can be shortened to `$install/` for more readability.
- `$opam/$OVERSION/lib/$NAME/` contains the libraries associated to the installed package `$NAME` with the compiler version `$OVERSION`. `$opam/$OVERSION/lib/` can be shortened to `$lib/` for more readability.
- `$opam/$OVERSION/doc/$NAME/` contains the documentation associated to the installed package `NAME` with the compiler version `$OVERSION`. `$opam/OVERSION/doc/` can be shortened to `$doc/` for more readability.
- `$opam/$OVERSION/bin/` contains the program files for all installed packages with the compiler version `$OVERSION`. `$opam/$OVERSION/bin/` can be shortened to `$bin/` for more readability.
- `$opam/$OVERSION/build/$NAME.$VERSION/` is a tempory folder used to build package `$NAME` with version `$VERSION`, with compiler version `$OVERSION`. `$opam/$OVERSION/build/` can be shortened to `$build/` for more readability.

- `$opam/$OVERSION/reinstall` contains the list of packages which has been changed upstream since the last upgrade. This can happen for instance when a packager uploads a new archive or fix the OPAM file for a specific package version. Every package appearing in this file will be reinstalled (or upgraded if a new version is available) during the next upgrade when the current version of the compiler is `$OVERSION`. The file format is similar to the one described in §1.2.1.

1.1.3 Repository-specific state

Configuration files for OPAM repositories `REPO` are stored in `$opam/repo/$REPO`. Repositories can be of different kinds (stored on the local filesystem, available via HTTP, available using a custom binary protocol, stored under git, ...); they all share the same base filesystem which is initialized using the `opam-<kind>-init` script (see §1.4.1) and use only the `opam-<kind>-update` (see §1.4.2) and `opam-<kind>-upload` (see §1.4.4) scripts to exchange data between the client and the corresponding OPAM repository.

- `$opam/repo/$REPO/kind` contains the kind associated to the OPAM repository `$REPO` have. The kind is stored as a single word (containing only letters and digits) and specifies which `opam-<kind>-*` scripts to call when updating and uploading this repository.
- `$opam/repo/$REPO/address` contains the address of the OPAM repository `$REPO`. This address is passed as argument to the `opam-<kind>-*` scripts.
- `$opam/repo/$REPO/opam/$NAME.$VERSION.opam` is the OPAM specification for the package `$NAME` with version `$VERSION` (which might not be installed). The format of OPAM files is described in §1.2.5.
- `$opam/repo/$REPO/descr/$NAME.$VERSION` contains the textual description for the version `$VERSION` of package `$NAME` (which might not be installed). The first line of this file is the package synopsis.
- `$opam/repo/$REPO/archives/$NAME.$VERSION.tar.gz` contains the source archives for the version `$VERSION` of package `$NAME`. This folder is populated by calling the corresponding `opam-<kind>-download` script (see §1.4.3).
- `$opam/repo/$REPO/updated` contains the new available packages which have not yet been synchronized with the client state. This file is created by the `opam-<kind>-update` script (see §1.4.2). If the file empty, this means that the client state is up-to-date. The file format is the same as the one described in §1.2.1.
- `$opam/repo/$REPO/upload/$NAME.$VERSION/` contains the OPAM, description and archive files to upload to the OPAM repository for the version `$VERSION` of package `$NAME`. The script `opam-<kind>-update` script (see §1.4.4) read the contents of `upload/` and send it to the repository (if the repository support upload).

1.2 File syntax

1.2.1 List of packages

The following configuration files: `$opam/$OVERSION/installed`, `$opam/$OVERSION/reinstall`, and `$opam/repo/$REPO/updated` follow a very simple syntax. The file is a list of lines which contains a space-separated name and a version. Each line `$NAME $VERSION` means that the version `$VERSION` of package `$NAME` has been compiled with OCaml version `$OVERSION` and has

been installed on the system in `$lib/$NAME` and `$bin/`.

For instance, if `batteries` version `1.0+beta` and `ocamlfind` version `1.2` are installed, then `$opam/$OVERSION/installed` will contain:

```
batteries 1.0+beta
ocamlfind 1.2
```

1.2.2 Index of packages

`$opam/repo/index` follows a very simple syntax: each line of the file contains a space separated list of words `$NAME $REPO` specifying that all the versions of package `$NAME` are available in the OPAM repository `$REPO`. The file contains information on all available packages (e.g. not only on the installed one).

For instance, if `batteries` version `1.0+beta` is available in the `testing` repository and `ocamlfind` version `1.2` is available in the `default` and `testing` repositories (where `default` is one being used), then `$opam/repo/index` will contain:

```
batteries testing
ocamlfind default
```

1.2.3 General syntax

Most of the files in the client and server states share the same syntax defined in this section.

Base types The base types for values are:

- **BOOL** is either `true` or `false`
- **STRING** is a doubly-quoted OCaml string, for instance: `"foo"`, `"foo-bar"`, ...
- **SYMBOL** contains only non-letter and non-digit characters, for instance: `=`, `<=`, ... Some symbols have a special meaning and thus are not valid **SYMBOLs**: `"() [] { } :"`.
- **IDENT** starts by a letter and is followed by any number of letters, digit and symbols, for instance: `foo`, `foo-bar`, ...

Compound types Types can be composed together to build more complex values:

- `X Y` is a space-separated pair of value.
- `X | Y` is a value of type either `X` or `Y`.
- `?X` is zero or one occurrence of a value of type `X`.
- `X+` is a space-separated list of values of at least one value of type `X`.
- `X*` is a space-separated list of values of values of type `X` (it might contain no value).

All structured OPAM files share the same syntax:

```

<file>  := <item>*

<item>  := IDENT : <value>
        | ?IDENT: <value>
        | IDENT STRING { <item>+ }

<value> := BOOL
        | STRING
        | SYMBOL
        | IDENT
        | [ <value>+ ]
        | value { <value>+ }

```

1.2.4 Global configuration file

\$opam/config follows the syntax defined in §1.2.3 with the following restrictions:

```

<file> :=
  opam-version: "1"
  ?repositories: [ <repo>+ ]
  ocaml-version: STRING

<repo> := STRING { STRING }

```

The field `repositories` contains the list of OPAM repositories with their kind.

The field `ocaml-version` corresponds to the current OCaml compiler (available in the path).

1.2.5 Package specification files: .opam

\$opam/opam/\$NAME.\$VERSION.opam follows the syntax defined in §1.2.3 with the following restrictions:

```

<file> :=
  opam-version: "1"
  package STRING {
    version:      STRING
    maintainer:   STRING
    ?subst:       [ STRING+ ]
    ?build:       [ command+ ]
    ?remove:      command
    ?depends:      <cnf-formula>
    ?optdeps:     <cnf-formula>
    ?conflicts:   <and-formula>
    ?libraries:   [ STRING+ ]
    ?syntax:      [ STRING+ ]
  }

<cnf-formula> := <cnf-formula> <cnf-formula>
               | ( <cnf-formula> )
               | <or-formula>

<or-formula> := STRING
               | STRING { <constraints> }
               | <or-formula> '||' <or-formula>

```

```

<and-formula> := STRING
               | STRING { <constraints> }
               | <and-formula> <and-formula>

<constraints> := <comp> STRING
               | <constraints> <constraints>

<comp>        := '=' | '<' | '>' | '>=' | '<='
<command>     := [ STRING+ ]

```

- The first line specifies the OPAM version.
- The string after **package** should not contain any dot ('.') nor space (' ').
- The contents of **version** is \$VERSION. The content of **maintainer** is the contact address of the package maintainer.
- The content of **subst** is the list of files to substitute variable (see §1.2.8 for the file format and §1.3.7 for the semantic of file substitution).
- The content of **build** is the list of commands to run in order to build the package libraries. The build script should build all the libraries and syntax extensions exported by the package and it should produce the platform-specific configuration and install files (e.g. \$NAME.config and \$NAME.install, see §1.2.6 and §1.2.7).
- The content of **remove** is the command to run before deleting the installed file.
- The **depends** and **conflicts** fields contain respectively *cnf*- and *and*-formulas over package names, optionally parametrized by version constraints. An expression is either:
 - A package name: "foo";
 - A package name with version constraints: "foo" (>= "1.2" & <= "3.4")
 - A disjunction of formulas: E | F
 - A conjunction of formulas: E F
 - A formula between parenthesis: (E)

For instance "foo" {<= "1.2"} ("bar" | "gna" {= "3.14"}) is a valid formula whose semantic is: *both any version of package "foo" lesser or equal to 1.2 and either any version of package "bar" or the version 3.14 of package "gna"*.

- The **optdeps** field contains a *cnf*-formula over package names, optionally parametrized by version constraints. When installing a package, OPAM will not try to install the optional dependencies. However, it will check if the package is an optional dependency of already installed packages. If it is the case, it will re-install the packages (and their transitive forward-dependency closure).
- The **libraries** and **syntax** fields contain the libraries and syntax extensions defined by the package.

1.2.6 Package configuration files: .config

`$opam/VERSION/config/NAME.config` follows the syntax defined in §1.2.3, with the following restrictions:

```
<file>      :=
  opam-version: "1"
  <item>*
<item>      := <def> | <section>
<section> :=
  <kind> STRING {
    ?asmcomp: [ STRING+ ]
    ?bytecomp: [ STRING+ ]
    ?asmlink : [ STRING+ ]
    ?bytelinek: [ STRING+ ]
    ?requires: [ STRING+ ]
  }
  <def>*
<kind>      := library | syntax
<def>       := IDENT: BOOL
              | IDENT: STRING
              | IDENT: [ STRING+ ]
```

`$NAME.config` contains platform-dependent information which can be useful for other libraries or syntax extensions that want to use libraries defined in the package `$NAME`.

Local and global variables The definitions “`IDENT: BOOL`”, “`IDENT: STRING`” and “`IDENT: [STRING+]`”, are used to defined variables associated to this package, and are used to substitute variables in template files (see §1.2.8):

- `%{$NAME:$VAR}%` will refer to the variable `$VAR` defined at the root of the configuration file `$config/NAME.config`.
- `%{$NAME.$LIB:$VAR}%` will refer to the variable `$VAR` defined in the `library` or `syntax` section named `$LIB` in the configuration file `$config/$NAME.config`.

Library and syntax sections Each `library` and `syntax` section defines an OCaml library and the specific compilation flags to enable when using and linking with this library.

The distinction between libraries and syntax extensions is only useful at compile time to know whether the options should be used as compilation or pre-processing arguments (ie. should they go on the compiler command line or should they be passed to the `-pp` option). This is the responsibility of the build tool to do the right thing and the `<kind>` of sections is only used for documentation purposes in OPAM.

The available options are:

- `asmcomp` are compilation options to give to the native compiler (when using the `-c` option)
- `bytecomp` are compilation options to give to the bytecode compiler (when using the `-c` option)
- `asmlink` are linking options to give to the native compiler
- `bytelinek` are linking options to give to the bytecode compiler

- **requires** is the list of libraries and syntax extensions the current block is depending on. The full list of compilation and linking options is built by looking at the transitive closure of dependencies. The contents of **deps** is the list of libraries or syntax extension the current section depends on. Note that we do not refer here to any package name, as multiple packages can expose libraries with the same name and interface and thus we want the user to be able to switch between them easily.

1.2.7 Package installation files: `.install`

`$opam/VERSION/install/NAME.install` follows the syntax defined in §1.2.3 with the following restrictions:

```
<file> :=
  opam-version: "1"
  ?lib:  [ STRING+ ]
  ?bin:  [ <mv>+ ]
  ?doc:  [ STRING+ ]
  ?misc: [ <mv>+ ]

<mv> := STRING
      | STRING { STRING }
```

- Files listed under `lib` are copied to `$lib/$NAME/`.
- Files listed under `bin` are copied to `$bin/` (they can be renamed using `$SRC { $DST }`; in this case `$SRC` should be a simple filename, ie. it should not start with a directory name).
- Files listed under `doc` are copied to `$doc/$NAME/`.
- Files listed under `misc` should be processed as follows: for each pair `$SRC { $DST }`, the tool should ask the user if he wants to install `$SRC` to the absolute path `$DST`.

1.2.8 Substitution files

All of the previous files can be generated using a special mode of `opam` which can perform tests and substitutes variables (see §1.3.7 for the exact command to run). Substitution files contains some templates which will be replaced by some contents. The syntax of templates is the following:

- templates such as `%{$NAME:$VAR}%` are replaced by the value of the variable `$VAR` defined at the root of the file `$config/NAME.config`.
- templates such as `%{$NAME.$LIB:$VAR}%` are replaced by the value of the variable `$VAR` defined in the `$LIB` section in the file `$config/PACKAGE.config`

1.3 Client commands

1.3.1 Creating a fresh client state

When an end-user starts OPAM for the first time, he needs to initialize `$opam/` in a consistent state. In order to do so, he should run:

```
$ opam init [-kind $KIND] $REPO $ADDRESS
```

Where:

- `$KIND` is the kind of OPAM repository (default is `http`);
- `$REPO` is the name of the repository (default is `default`); and
- `ADDRESS` is the repository address (default is `http://opam.ocamlpro.com/pub`).

This command will:

1. Create the file `$opam/config` (as specified in §1.2.4)
2. Create an empty `$opam/$OVERSION/installed` file, `$OVERSION` is the version from the OCaml used to compile `$opam`. In particular, we will not fail now if there is no `ocamlc` in `$path`.
3. Initialize `$opam/repo/$REPO` by running `'opam-$KIND-init $ADDRESS'` (see §1.4.2). If the script cannot be found in the path, the command should be canceled and should return a well-defined error.
4. Symlink all OPAM and description files (ie. create a symbolic link from every file in `$opam/repo/$REPO/opam/` to `$opam/opam/` and from every file in `$opam/repo/$REPO/descr/` to `$opam/descr/`).
5. Create `$opam/repo/index` and for each version `$VERSION` of package `$NAME` appearing in the repository, append the line `'$REPO $NAME $VERSION'` to the file.
6. Create the empty directories `$opam/archives`, `$lib/`, `$bin/` and `$doc/`.

1.3.2 Listing packages

When an end-user wants to have information on all available packages, he should run:

```
$ opam list
```

This command will parse `$opam/$OVERSION/installed` to know the installed packages, and `$opam/opam/*.opam` to get all the available packages. It will then build a summary of each packages. The description of each package will be read in `$opam/descr/` if it exists.

For instance, if `batteries` version 1.1.3 is installed, `ounit` version 2.3+dev is installed and `camomille` is not installed, then running the previous command should display:

```
batteries    1.1.3  Batteries is a standard library replacement
ounit        2.3+dev Test framework
camomille     --    Unicode support
```

1.3.3 Getting package info

In case the end-user wants a more details view of a specific package, he should run:

```
$ opam info $NAME
```

This command will parse `$opam/$OVERSION/installed` to get the installed version of `$NAME`, will process `$opam/repo/index` to get the repository where the package comes from and will look for `$opam/opam/$NAME.*.opam` to get available versions of `$NAME`. It can then display:

```

package: $NAME
version: $VERSION                # '--' if not installed
versions: $VERSION1, $VERSION2, ...
libraries: $LIB1, $LIB2, ...
syntax: $SYNTAX1, $SYNTAX2, ...
repository: $REPO
description:
    $SYNOPSIS

    $LINE1
    $LINE2
    $LINE3
    ...

```

1.3.4 Installing a package

When an end-user wants to install a new package, he should run:

```
$ opam install $NAME
```

This command will:

1. Compute the transitive closure of dependencies and conflicts of packages using the dependency solver (see §1.3.11). If the dependency solver returns more than one answer, the tool will ask the user to pick one, otherwise it will proceed directly. The dependency solver should also mark the packages to recompile.
2. The dependency solver sorts the collections of packages in topological order. Then, for each of them do:
 - (a) Check whether the package is already installed by looking for the line `$NAME $VERSION` in `$opam/$VERSION/installed`. If not, then:
 - (b) Look into the archive cache to see whether it has already been downloaded. The cache location is: `$opam/archives/$NAME.VERSION.tar.gz`
 - (c) If not, process `$opam/repo/index/` to get the repository `$REPO` where the archive is available, get the repository kind by looking at `$opam/repo/$REPO/kind` and then ask the repository to download the archive by calling `opam-$KIND-download` (see §1.4.3).
Once this is done, symlink the archive in `$opam/archives`.
 - (d) Decompress the archive into `$build/$NAME.$VERSION/`.
 - (e) Substitute the required files.
 - (f) Run the list of commands to build the package with `$bin` in the path.
 - (g) Process `$build/$NAME.$VERSION/$NAME.install` to install the created files. The file format is described in §1.2.7.
 - (h) Install the installation file `$build/$NAME.$VERSION/$NAME.install` in `$install/` and the configuration file `$build/$NAME.$VERSION/$NAME.config` in `$config/`.

1.3.5 Updating index files

When an end-user wants to know what are the latest packages available, he will write:

```
$ opam update
```

This command will follow the following steps:

- For each repositories in `$opam/config`, run the appropriate `opam-$KIND-update` script (see §1.4.2).
- For each repositories in `$opam/config`, process `$opam/repo/$REPO/updated` and update `$opam/repo/index`, `$opam/opam/` and `$opam/desc` accordingly (ie. add the right lines in `$opam/repo/index` and create the missing symlinks). Here, the order in which the repositories are specified is important: the first repository containing a given version for a package will be the one providing it (this can be changed manually by editing `$opam/repo/index` later).
- For each line `$REPO $NAME $VERSION` in `$opam/repo/index`, if the version `$VERSION` of package `$NAME` has been modified upstream (ie. if the line `$NAME $VERSION` appears in `$opam/repo/$REPO/$updated`) and if the package is already installed (ie. it appears in `opam/$OVERSION/installed`), then update `$opam/$OVERSION/reinstall` accordingly (for each compiler version `$OVERSION`).

Packages in `$opam/$OVERSION/reinstall` will be reinstalled (or upgraded if a new version is available) on the next `opam upgrade` (see §1.3.6), with `$OVERSION` being the current compiler version when the upgrade command is run.

- Delete each `$opam/repo/$REPO/$updated`

1.3.6 Upgrading installed packages

When an end-user wants to upgrade the packages installed on his host, he will write:

```
$ opam upgrade
```

This command will:

- Call the dependency solver (see §1.3.11) to find a consistent state where **most** of the installed packages are upgraded to their latest version. Moreover, packages listed in `$opam/$OVERSION/reinstall` will be reinstalled (or upgraded if a new version is available). It will install each non-installed packages in topological order, similar to what it is done during the install step, See §1.3.4.
- Once this is done the command will delete `$opam/$OVERSION/reinstall`.

1.3.7 Getting package configuration

The first version of OPAM contains the minimal information to be able to use installed libraries. In order to do so, the end-user (or the packager) should run:

```

$ opam config -list-vars
$ opam config -var $NAME:$VAR
$ opam config -var $NAME.$LIB:$VAR
$ opam config -subst $FILENAME+
$ opam config [-r] -I          $NAME+
$ opam config [-r] -bytecomp $NAME.$LIB+
$ opam config [-r] -asmcomp  $NAME.$LIB+
$ opam config [-r] -bytelink $NAME.$LIB+
$ opam config [-r] -asmlink  $NAME.$LIB+

```

- `-list-vars` will return the list of all variables defined in installed packages (see §1.2.6)
- `-var $var` will return the value associated to the variable `$var`
- `-subst $FILENAME` replace any occurrence of `%{$NAME:$VAR}%` and `%{$NAME.$LIB:$VAR}%` as specified in §1.2.8 in `$FILENAME.in` to create `$FILENAME`.
- `-I $NAME` will return the list of paths to include when compiling a project using the package `$NAME` (`-r` gives a result taking into account the transitive closure of dependencies).
- `-bytecomp`, `-asmcomp`, `-bytelink` and `-asmlink` return the associated value for the section `$LIB` in the file `$config/$NAME.config` (`-r` gives a result taking into account the transitive closure of all dependencies).

1.3.8 Uploading packages

When a packager wants to create a package, he should:

1. create `$package/$NAME.$VERSION.opam` containing in the format specified in §1.2.5.
2. create a file describing the package
3. make sure the build scripts:
 - build the libraries and packages advertised in `$package/$NAME.$VERSION.opam`
 - generates a valid `$package/$NAME.install` containing the list of files to install (the file format is described in 1.2.7).
 - generates a valid `$package/$NAME.config` containing the configuration flags for libraries exported by this package (the file format is described in 1.2.6).
4. create an archive `$NAME.$VERSION.tar.gz` with the sources he wants to distribute.
5. run the following command:

```
$ opam upload -opam $OPAM -descr $DESCR -archive $ARCHIVE -repo $REPO
```

This command will parse `$OPAM` to get the package name and version and it will:

- move `$OPAM` to `$opam/repo/$REPO/upload/$NAME.$VERSION.opam`
- move `$DESCR` to `$opam/repo/$REPO/descr/$NAME.$VERSION`
- move `$ARCHIVE` to `$opam/repo/$REPO/archives/$NAME.$VERSION.tar.gz`

It will then call `opam- $\$KIND$ -upload` to upload the files upstream (see §1.4.4).

This command will work only for writable repositories; in case of errors (for instance if the repository is read-only) the script returns a non-zero exit code and stores the error message in `$opam/repo/ $\$REPO$ /error` (see §1.4.5).

1.3.9 Removing packages

When the user wants to remove a package, he should write:

```
$ opam remove $NAME
```

This command will check whether the package `$NAME` is installed, and if yes, it will display to the user the list packages that will be uninstalled (ie. the transitive closure of all forward-dependencies). If the user accepts the list, all the packages should be uninstalled, and the client state should be let in a consistent state.

1.3.10 Managing OPAM repository

When the user wants to manage OPAM repositories, he should write:

```
$ opam remote -list
$ opam remote -rm $REPO
$ opam remote -add [-kind $KIND] $REPO $ADDRESS
```

- `-list` lists the current repositories by looking at `$opam/config`
- `-rm $REPO` deletes `$opam/repo/$REPO` and removes `$REPO` from the `repositories` list in `$opam/config`. Then, for each package in `$opam/repo/index` it updates the link between packages and repositories (ie. it either deletes packages or symlink them to the new repository containing the package).
- `-add [-kind $KIND] $REPO $ADDRESS` initializes `$REPO` as described in §1.3.1.

1.3.11 Dependency solver

Dependency solving is a hard problem and we do not plan to start from scratch implementing a new SAT solver. Thus our plan to integrate (as a library) the Debian dependency solver for CUDF files, which is written in OCaml.

- the dependency solver should run on the client; and
- the dependency solver should take as input a list of packages (with some optional version information) the user wants to install, upgrade and remove and it should return a consistent list of packages (with version numbers) to install, upgrade, recompile and remove.

1.4 OPAM repository scripts

As stated above, OPAM repositories can be of different kinds. A given kind `$KIND` has associated scripts `opam- $\$KIND$ -init`, `opam- $\$KIND$ -update`, `opam- $\$KIND$ -download` and `opam- $\$KIND$ -upload` which should behave in the way specified below.

1.4.1 Init script

```
$ opam-$KIND-init $ADDRESS
```

This script should contact the OPAM repository located at address `$ADDRESS` and populate its local state (ie. `./opam/`, `./descr/` and `archives`) as specified in §1.1.3.

1.4.2 Update script

```
$ opam-$KIND-update $ADDRESS
```

This script should get the list of newly available packages, by contacting the OPAM repository located at address `$ADDRESS` and update `./opam` and `./descr`.

It should then create (or update) `./updated` to reflect the newly available packages and the one who changed upstream..

1.4.3 Download script

```
$ opam-$KIND-download $ADDRESS $NAME.$VERSION
```

This script should get `$NAME.$VERSION.tar.gz` from the OPAM repository located at address `ADDRESS` and store it into `$opam/repo/$REPO/archives/` (this should overwrite any previously existing file if one is already there).

1.4.4 Upload script

```
$ opam-$KIND-init $opam/repo/$REPO
```

This script should send the contents of `$opam/repo/$REPO/upload/` to the repository whose address is stored in `$opam/repo/$REPO/address` and remove the folder contents when it's done.

1.4.5 Error handling

In case any of the above scripts return with a non-zero exit code, the full error message will be stored into `opam/repo/$REPO/error`

1.4.6 RSYNC repository

The first Milestone includes the necessary scripts to support `rsync` repositories. In order to do so:

- The OPAM repository is a filesystem similar to what is described in `$opam/repo/$REPO`.
- `opam-rsync-init $PATH $ADDRESS` simply runs something like:

```
$ cd $PATH && rsync [...] $ADDRESS/opam && rsync [...] $ADDRESS/descr
```

- `opam-rsync-download $PATH $NAME $VERSION` will simply run something like:

```
$ cd $PATH && rsync [...] $ADDRESS/archives/$NAME.$VERSION.tar.gz
```

- `opam-rsync-update $PATH` will also run `rsync` and filter the command output to populate `$opam/repo/$REPO/updated` accordingly.

- `opam-rsync-upload $PATH` will return an error if the OPAM repository is not writable (for instance if the repository is an HTTP server). In this case the server state can be updated manually by copying the new files in the right place on the server filesystem.

2 Milestone 2: OPAM server

The only OPAM repository kind available in Milestone 1 is `rsync`.

This second milestone describes a new kind of repository: `server`. The OPAM server uses a custom binary protocol between a client (the `opam-server-*` scripts) and a server (`opam-server` which is run on the server). This is intended to be the preferred way to upload new packages to OPAM repositories as it does not require any external authentication mechanisms.

2.1 Server state

In this section, `$opamserver` refers to the filesystem subtree containing the server state. Default directory is `$home/.opam-server`. The state is similar to what is described in §1.1.1:

- `$opamserver/opam/$NAME.$VERSION.opam` is the OPAM specification for the package `$NAME` with version `$VERSION` (which might not be installed). The format of OPAM files is described in §1.2.5.
- `$opamserver/descr/$NAME.$VERSION` contains the description for the version `$VERSION` of package `$NAME` (which might not be installed). The first line of this file is the package synopsis.
- `$opamserver/archives/$NAME.$VERSION.tar.gz` contains the source archives for the version `$VERSION` of package `$NAME`.

2.2 Binary Protocol

The protocol is very simple, there is one kind of messages corresponding to each basic client actions (which will be used by the `opam-server-*` scripts). The only “complex” parts reside in the basic authentication mechanisms: the server stores a key for each package (every version shares the same key); the key is sent to the client the first time a new package is uploaded and key consistency is checked each time a new version of the package is uploaded.

- Communication between clients and servers always starts by a hand-shake to agree on the protocol version.
- All the basic values (names, versions and binary data) are represented as OCaml strings and are marshaled as a 64-bits integer (the string size) followed by the string contents.
- Messages are marshaled using a simple binary protocol: the first byte represents the message number, and then each string argument is stacked sequentially. The list of messages *from the client to server* is:

#	Client-to-Server Message	Contents	Description
0	ClientVersion	version: string	Send the client version to the server
1	GetList	–	Ask for the list of all available OPAM files
2	GetOPAM	name : string version: string	Ask for the binary representation of a given OPAM file
3	GetDescr	name : string version: string	Ask for the binary representation of a given description file
4	GetArchive	name : string version: string	Ask for the binary representation of a given archive file
5	NewPackage	name : string version: string opam : string descr : string archive: string	Create a new package on the server. The client should provide the OPAM and descr files and the source archive.
6	NewVersion	name : string version: string opam : string descr : string archive: string key : string	Upload a new version of an already existing package on the server. The client should also provide a security key

- Answers from the server are encoded in the same way (ie, a byte for the message number, followed by the sequential marshaling of the arguments). Pairs are stacked sequentially. List are encoded by stacking first the list length in a 64-bit integer and then all the elements of the list in sequential order. The list of messages *from servers to clients* is:

#	Server-to-Client Message	Contents	Description
0	ServerVersion	version: string	Return the server version
1	PackageList	list : (string*string) list	Return the list of available package names and versions
2	OPAM	opam : string	Return an OPAM file
3	Descr	descr : string	Return an OPAM file
4	Archive	archive: string	Return an archive file
5	Key	key : string	Return a security key
6	OK	–	The update went OK
7	Error	error : string	An error occurred

Note that when an error is raised by an arbitrary function at server side, the client receives **Error**.

2.3 Server authentication

The server should be able to ask for basic credential proofs. The protocol can be sketched as follows:

- packagers store keys in `$opam/keys/NAME`. These keys are random strings of size 128.
- the server stores key hashes in `$opamserver/hashes/NAME`.
- when a packager uploads a fresh package the server returns a random key. Clients then stores that key in `$opam/keys/NAME`.
- when a packager wants to uplaod a new version of an existing package, the client also sends the key. The server then checks whether the hash of the key is equal to the one

stored in `$opamserver/ashes/NAME`; if yes, it updates the package and return OK, if no if it returns an Error.

2.4 Scripts

The initialization script: `opam-server-init`, the update script: `opam-server-update`, the download script: `opam-server-download` and the upload script: `opam-server-upload` will work as expected, using the binary protocol defined in §2.2 in order to correctly update the client state as defined in §1.4.

3 Milestone 3: Multiple compiler versions

This milestone focus on the support of multiple compiler versions.

3.1 Compiler Description Files

For each compiler version `OVERSION`, the client state will be extended with the following files:

- `$opam/compilers/OVERSION.comp`

The syntax of `.comp` files follows the one described in §1.2.3 with the following restrictions:

```
<file> :=
  opam-version: "1"
  name:        STRING
  src:         STRING
  make:        [ STRING+ ]
  ?patches:    [ STRING+ ]
  ?configure:  [ STRING+ ]
  ?bytecomp:   [ STRING+ ]
  ?asmcomp:    [ STRING+ ]
  ?bytelink:   [ STRING+ ]
  ?asmlink:    [ STRING+ ]
  ?packages:   [ STRING+ ]
  ?requires:   [ STRING+ ]
  ?pp:         [ <ppflag>+ ]
  ?preinstalled: BOOL

<ppflag> := CAMLP4 { STRING+ }
           | STRING+
```

- `name` is the compiler name, it should be identical to the filename.
- `src` is the location where this version can be downloaded. It can be:
 - an archive available in the local filesystem
 - an archive available via `http` or `ftp`
 - a version-controlled repository under `svn` or `git` (with the expectation that these tools are installed on the user host).

- **patches** are optional patch addresses, available via **http**, **ftp** or locally on the filesystem.
- **configure** are the optional flags to pass to the configure script. The order is relevant: **-prefix=\$opam/OVERSION/** will be automatically added at the end to these options. Remark that if these flags contain **-bindir**, **-libdir**, and **-mandir**, then every **-prefix** will be ignored by **configure**.
- **make** are the flags to pass to **make**. It must at least contain some target like **world** or **world.opt**.
- **bytecomp**, **asmcomp**, **bytelinek** and **asmlink** are the compilation and linking flags to pass to the OCaml compiler. They will be taken into account by the **opam config** command (see §1.3.7).
- **packages** is the list of packages to install just after the compiler installation finished. These libraries will not consider what is in the **requires** nor **pp** (as **requires** and **pp** might want to use things already installed with **packages**).
- **requires** is a list of libraries and syntax extensions dependencies which will be added to every packages installed with this compiler. The libraries and syntax extensions should be present in packages defined in **packages**, otherwise an error should be thrown.
- **pp** is the command to use with the **-pp** command-line argument. It is either a full command line or a **camlp4** command, such as **CAML4 ["pp-trace"]**: this will look for the compilation flags for the syntax extension **"pp-trace"** and expand the **camlp4** command-line accordingly. All the syntax extensions used should be present in **packages**.
- **preinstall** is **true** when the version of the compiler available in the path is the same as **name**.

For instance the file, **3.12.1+memprof.comp** describes OCaml, version 3.12.1 with the memory profiling patch enabled:

```
opam-version: "1"
name:        "3.12.1"
src:         "http://caml.inria.fr/pub/distrib/ocaml-3.12/ocaml-3.12.1.tar.gz"
make:        [ "world" "world.opt" ]
patches:     [ "http://bozman.cagdas.free.fr/documents/ocamlmemprof-3.12.0.patch" ]
```

And the file **trunk-g-notk-byte.comp** describes OCaml from SVN trunk, with no *tk* support and only in bytecode, and all the libraries built with **-g**:

```
opam-version: "1"
name:        "trunk-g-notk-byte"
src:         "http://caml.inria.fr/pub/distrib/ocaml-3.12/ocaml-3.12.1.tar.gz"
configure:   [ "-no-tk" ]
make:        [ "world" ]
bytecomp:    [ "-g" ]
bytelinek:   [ "-g" ]
```

3.2 Switching compiler version

If the user wants to switch to an other compiler version, he should run:

```
$ opam switch [-clone] [-alias $ALIAS] $OVERSION
```

This command will:

- If `$ALIAS` is not set, set it to `$OVERSION`
- Look for an existing `$opam/$ALIAS` directory.
 - If it exists, then change the `ocaml-version` content to `$ALIAS` in `$opam/config`.
 - If it does not exist, look for an existing `$opam/compilers/OVERSION.comp`. If the file does not exists, the command will fail with a well-defined error.
 - If the file exist, then build the new compiler with the right options (and pass `--prefix $opam/$ALIAS` to `./configure`) and initialize everything in `$opam/` in a consistent state as if “`opam init`” has just been called.
 - Update the file `$opam/aliases` with the line `$ALIAS $OVERSION`
- If the `-clone` option is set, the command will try to install the packages that were installed before switching (that are not currently installed). In case the new version contains installed packages that were not installed before switching, it will try to keep them.

In short, the heuristic is to install the maximum of previous packages and remove the minimum. The success depends on the compatibility of the existing packages with respect to this new `$OVERSION`.

3.3 OPAM repositories

The state of OPAM repositories is extended with the directory `$opam/repo/$repo/compiler` containing the compiler description files. When a repository is updated, this directory is updated as well. It is advised to no have two different compiler description files with the same name in two different repositories.

The binary protocol between client and OPAM server is also updated to deal with compiler description updates.

4 Milestone 4: OPAM and git

4.1 OPAM repository under git

4.2 Package repository under git

5 Milestone 9: Parallel Build

A new field, `cores` is added to `$opam/config`. It contains the number of parallel processes that OPAM will spawn to compile the packages.

6 Milestone 8: Version Pinning